

```
import sys
!{sys.executable} -m pip install -U pandas-profiling[notebook]
!jupyter nbextension enable --py widgetsnbextension
!pip install matplotlib
!pip install graphviz
```

```
Requirement already satisfied: pandas-profiling[notebook] in /usr/local/lib/python3.7/dist-packages (1.4.1)
WARNING: pandas-profiling 1.4.1 does not provide the extra 'notebook'
Requirement already satisfied: jinja2>=2.8 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (2.11.3)
Requirement already satisfied: six>=1.9 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (1.16.0)
Requirement already satisfied: matplotlib>=1.4 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (3.5.3)
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (1.1.5)
Requirement already satisfied: MarkupSafe>=0.23 in /usr/local/lib/python3.7/dist-packages (from jinja2>=2.8->pandas-profiling[notebook]) (2.0.1)
Requirement already satisfied: cyclizer>=0.10 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (0.10.0)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=1.4->pandas-profiling[notebook]) (1.3.1)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (1.19.5)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->pandas>=0.19->pandas-profiling[notebook]) (3.0.7)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->pandas>=0.19->pandas-profiling[notebook]) (2021.3)
Enabling notebook extension jupyter-js-widgets/extension...
- Validating: OK
Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages (3.5.3)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cyclizer>=0.10 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (0.10.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->pandas>=0.19->pandas-profiling[notebook]) (3.0.7)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (1.19.5)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (1.16.0)
Requirement already satisfied: graphviz in /usr/local/lib/python3.7/dist-packages (0.10)
```

```
from google.colab import files
uploaded = files.upload()
```

```
import pandas as pd
from scipy.io import arff
import numpy as np
```

```
data_file = "stroke_preprocessed.arff"
data = arff.loadarff(data_file)
```

```
df = pd.DataFrame(data[0])
for col in df.columns:
    if df[col].dtype == 'object':
        # Ensure data isn't read as bytes but rather as strings from file
```

```
df[col] = df[col].str.decode('utf-8')
# Examine data types
print(df.dtypes)
```

```
"id"                float64
"gender"            object
"age"              float64
"hypertension"      object
"heart_disease"     object
"ever_married"      object
"work_type"         object
"residence_type"    object
"avg_glucose_level" float64
"bmi"              float64
"smoking_status"    object
"stroke"            object
dtype: object
```

```
# Display first 10 rows
df.head(10)
```

```
# Examine meta info about data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
```

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	"id"	5110 non-null	float64
1	"gender"	5110 non-null	object
2	"age"	5110 non-null	float64
3	"hypertension"	5110 non-null	object
4	"heart_disease"	5110 non-null	object
5	"ever_married"	5110 non-null	object
6	"work_type"	5110 non-null	object
7	"residence_type"	5110 non-null	object
8	"avg_glucose_level"	5110 non-null	float64
9	"bmi"	5110 non-null	float64
10	"smoking_status"	5110 non-null	object
11	"stroke"	5110 non-null	object

dtypes: float64(4), object(8)

memory usage: 479.2+ KB

# The original 201 null values were all from bmi column, and they have been replaced by place

# Convert the 5000 values back into null values

df = df.replace(5000.0, np.nan)

# Check head of dataset again

df.head(10)

# Check structure of data types to ensure bmi remains float

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   "id"                   5110 non-null   float64
1   "gender"               5110 non-null   object
2   "age"                  5110 non-null   float64
3   "hypertension"         5110 non-null   object
4   "heart_disease"        5110 non-null   object
5   "ever_married"         5110 non-null   object
6   "work_type"            5110 non-null   object
7   "residence_type"       5110 non-null   object
8   "avg_glucose_level"    5110 non-null   float64
9   "bmi"                  4909 non-null   float64
10  "smoking_status"       5110 non-null   object
11  "stroke"               5110 non-null   object
dtypes: float64(4), object(8)
memory usage: 479.2+ KB
```

```
# Remove records with NAs from dataset
df_noNA = df
df_noNA = df_noNA.dropna()
df_noNA.head(10)
```

```
# Modify 'stroke' column so that 0 is coded as "No Stroke" and 1 is coded as "Stroke"
df_noNA.loc[df_noNA["stroke"] == '0', "stroke"] = "No Stroke"
```

```
df_noNA.loc[df_noNA['"stroke"' ] == '1', '"stroke"' ] = "Stroke"  
df_noNA.head(10)
```

```
# See if there are any extreme values in numeric data  
df_noNA.describe()
```

```
# Create list of categorical columns
cat_cols = ["gender", "hypertension", "heart_disease", "ever_married", "work_type",

# Create copy of a data frame in memory w/ a different name
df_dummy = df_noNA.copy()
# Convert only categorical feature into dummy/one-hot features
df_dummy = pd.get_dummies(df_noNA, columns = cat_cols, prefix = cat_cols)
# Print dataset
df_dummy
```

```
# Create train test set split
from sklearn.model_selection import train_test_split
```

```
# Set class name as "stroke", all else will be used as features
class_col_name = "stroke"
dummy_feature_names = df_dummy.columns[df_dummy.columns != class_col_name]
```

```

# 70% training, 30% test set split
x_train, x_test, y_train, y_test = train_test_split(df_dummy.loc[:, dummy_feature_names], df_

# Naive Bayes modeling
from sklearn.naive_bayes import MultinomialNB

# Create Multinomial NB Classifier
nb = MultinomialNB()

# Train model using training sets
nb.fit(x_train, y_train)

MultinomialNB()

# Predict response for test dataset
y_pred = nb.predict(x_test)

# Print Naive Bayes output
print("Number of features used: ", nb.n_features_)
print("Classes: ", nb.classes_)
print("Number of records for classes: ", nb.class_count_)
print("Log prior probability for classes: ", nb.class_log_prior_)
print("Log conditional probability for each feature given a class: ", nb.feature_log_prob_)

Number of features used: 24
Classes: ['No Stroke' 'Stroke']
Number of records for classes: [3291. 145.]
Log prior probability for classes: [-0.04311653 -3.16532954]
Log conditional probability for each feature given a class: [[-4.86773845e-03 -6.795029
-1.10563808e+01 -1.14237629e+01 -1.79362509e+01 -1.06200350e+01
-1.29804239e+01 -1.05742404e+01 -1.36665535e+01 -1.15561284e+01
-1.09735341e+01 -1.25702749e+01 -1.57390263e+01 -1.10803154e+01
-1.24452492e+01 -1.24852125e+01 -1.12629529e+01 -1.11845648e+01
-1.17028211e+01 -1.23527546e+01 -1.15025073e+01 -1.24167920e+01]
[-6.14395678e-03 -6.36670554e+00 -5.66931236e+00 -7.16182860e+00
-1.11140899e+01 -1.14296068e+01 -1.55567412e+01 -1.09027808e+01
-1.18190715e+01 -1.07776177e+01 -1.22245366e+01 -1.28486910e+01
-1.06739392e+01 -1.24656987e+01 -1.55567412e+01 -1.10458817e+01
-1.20013931e+01 -1.55567412e+01 -1.13670864e+01 -1.11622920e+01
-1.25610089e+01 -1.17955410e+01 -1.15494080e+01 -1.21227540e+01]]
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:103: FutureWarning:
warnings.warn(msg, category=FutureWarning)

# Confusion matrix and Evaluation metrics
from sklearn.metrics import confusion_matrix
cf = confusion_matrix(y_test, y_pred)

```

```

print("Confusion Matrix")
print(cf)
tn, fp, fn, tp = cf.ravel()
print("TP: ", tp, ", FP: ", fp, ", TN: ", tn, ", FN: ", fn)

```

```

Confusion Matrix
[[872 537]
 [ 21  43]]
TP:  43 , FP:  537 , TN:  872 , FN:  21

```

```

# Classifier report
from sklearn.metrics import classification_report
from sklearn import metrics

```

```

print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
No Stroke	0.98	0.62	0.76	1409
Stroke	0.07	0.67	0.13	64
accuracy			0.62	1473
macro avg	0.53	0.65	0.45	1473
weighted avg	0.94	0.62	0.73	1473

```

# Decision tree on dummy encoded data
from sklearn import tree
clf = tree.DecisionTreeClassifier(max_depth = 5) # 5 levels set
clf = clf.fit(x_train, y_train)

```

```

import graphviz
# Obtain unique class values to show on tree
class_values = df_dummy[class_col_name].unique()
print("class names: ", class_values)
dot_data = tree.export_graphviz(clf, out_file = None, feature_names = dummy_feature_names, cl

```

```

class names:  ['Stroke' 'No Stroke']

```

```

# Draw graph
graph = graphviz.Source(dot_data, format = "png")
# Graph won't display

```

```

# Perform prediction on test set
y_pred = clf.predict(x_test)

```

```

# Get decision tree report
from sklearn.metrics import classification_report
from sklearn import metrics
print(classification_report(y_test, y_pred))

```



	precision	recall	f1-score	support
No Stroke	0.96	0.99	0.97	1409
Stroke	0.12	0.03	0.05	64
accuracy			0.95	1473
macro avg	0.54	0.51	0.51	1473
weighted avg	0.92	0.95	0.93	1473

✓ 0s completed at 9:50 PM

