

```
# Balancing dataset using SMOTE oversampling method
# Then applied Logistic Regression, Naive Bayes Classifier, and Decision Tree algorithms
# Compared changes in accuracy rates when 10-fold cross-validation applied
```

```
import sys
```

```
!{sys.executable} -m pip install -U pandas-profiling[notebook]
```

```
!jupyter nbextension enable --py widgetsnbextension
```

```
!pip install matplotlib
```

```
!pip install graphviz
```

```
Requirement already satisfied: pandas-profiling[notebook] in /usr/local/lib/python3.7/dist-packages (1.4.1)
WARNING: pandas-profiling 1.4.1 does not provide the extra 'notebook'
```

```
Requirement already satisfied: Jinja2>=2.8 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (2.11.3)
```

```
Requirement already satisfied: six>=1.9 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (1.16.0)
```

```
Requirement already satisfied: matplotlib>=1.4 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (3.3.4)
```

```
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.7/dist-packages (from pandas-profiling[notebook]) (1.1.5)
```

```
Requirement already satisfied: MarkupSafe>=0.23 in /usr/local/lib/python3.7/dist-packages (from Jinja2>=2.8->pandas-profiling[notebook]) (2.0.1)
```

```
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=1.4->pandas-profiling[notebook]) (0.11.0)
```

```
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (1.21.0)
```

```
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=1.4->pandas-profiling[notebook]) (3.0.7)
```

```
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=1.4->pandas-profiling[notebook]) (1.3.2)
```

```
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (2.8.2)
```

```
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (4.1.1)
```

```
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pandas-profiling[notebook]) (2022.1)
```

```
Enabling notebook extension jupyter-js-widgets/extension...
```

```
Paths used for configuration of notebook:
```

```
  /root/.jupyter/nbconfig/notebook.json
```

```
- Validating: OK
```

```
Paths used for configuration of notebook:
```

```
  /root/.jupyter/nbconfig/notebook.json
```

```
Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages (3.3.4)
```

```
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib) (0.11.0)
```

```
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib) (2.8.2)
```

```
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib) (3.0.7)
```

```
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.7/dist-packages (from matplotlib) (1.21.0)
```

```
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib) (1.3.2)
```

```
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from matplotlib) (4.1.1)
```

```
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->matplotlib) (1.16.0)
```

```
Requirement already satisfied: graphviz in /usr/local/lib/python3.7/dist-packages (0.10)
```

```
from google.colab import files
```

```
uploaded = files.upload()
```

Choose files stroke\_preprocessed.arff

- **stroke\_preprocessed.arff**(n/a) - 368681 bytes, last modified: 06/03/2022 - 100% done

Saving stroke\_preprocessed.arff to stroke\_preprocessed.arff

```
import pandas as pd
```

```
from scipy.io import arff
```

```

import numpy as np

# Timing how long predictors take to run for efficiency calculations
# Import libraries
import time

data_file = "stroke_preprocessed.arff"
data = arff.loadarff(data_file)

df = pd.DataFrame(data[0])
for col in df.columns:
    if df[col].dtype == 'object':
        # Ensure data isn't read as bytes but rather as strings from file
        df[col] = df[col].str.decode('utf-8')
# Examine data types
print(df.dtypes)

    "id"                float64
    "gender"            object
    "age"               float64
    "hypertension"      object
    "heart_disease"     object
    "ever_married"      object
    "work_type"         object
    "residence_type"    object
    "avg_glucose_level" float64
    "bmi"               float64
    "smoking_status"    object
    "stroke"            object
    dtype: object

# Display first 10 rows
df.head(10)

```

	"id"	"gender"	"age"	"hypertension"	"heart_disease"	"ever_married"	"work_type"
0	9046.0	Male	67.0	0	1	Yes	Private
1	51676.0	Female	61.0	0	0	Yes	Self-employed
2	31112.0	Male	80.0	0	1	Yes	Private

```
# Examine meta info about data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   "id"                  5110 non-null   float64
1   "gender"              5110 non-null   object
2   "age"                 5110 non-null   float64
3   "hypertension"        5110 non-null   object
4   "heart_disease"       5110 non-null   object
5   "ever_married"        5110 non-null   object
6   "work_type"           5110 non-null   object
7   "residence_type"      5110 non-null   object
8   "avg_glucose_level"   5110 non-null   float64
9   "bmi"                 5110 non-null   float64
10  "smoking_status"      5110 non-null   object
11  "stroke"              5110 non-null   object
dtypes: float64(4), object(8)
memory usage: 479.2+ KB
```

```
# The original 201 null values were all from bmi column, and they have been replaced by place
# Convert the 5000 values back into null values
df = df.replace(5000.0, np.nan)
```

```
# Check head of dataset again
df.head(10)
```

	"id"	"gender"	"age"	"hypertension"	"heart_disease"	"ever_married"	"work_type"
0	9046.0	Male	67.0	0	1	Yes	Private
1	51676.0	Female	61.0	0	0	Yes	Self-employed
2	31112.0	Male	80.0	0	1	Yes	Private
3	60182.0	Female	49.0	0	0	Yes	Private
4	1665.0	Female	70.0	1	0	Yes	Self-employed

```
# Check structure of data types to ensure bmi remains float
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   "id"                        5110 non-null   float64
1   "gender"                    5110 non-null   object
2   "age"                        5110 non-null   float64
3   "hypertension"              5110 non-null   object
4   "heart_disease"             5110 non-null   object
5   "ever_married"              5110 non-null   object
6   "work_type"                  5110 non-null   object
7   "residence_type"            5110 non-null   object
8   "avg_glucose_level"         5110 non-null   float64
9   "bmi"                       4909 non-null   float64
10  "smoking_status"            5110 non-null   object
11  "stroke"                     5110 non-null   object
dtypes: float64(4), object(8)
memory usage: 479.2+ KB
```

```
# Remove records with NAs from dataset
df_noNA = df
df_noNA = df_noNA.dropna()
df_noNA.head(10)
```

	"id"	"gender"	"age"	"hypertension"	"heart_disease"	"ever_married"	"work_type"
0	9046.0	Male	67.0	0	1	Yes	Priv:
2	31112.0	Male	80.0	0	1	Yes	Priv:
3	60182.0	Female	49.0	0	0	Yes	Priv:
4	1665.0	Female	79.0	1	0	Yes	Si employ
5	56660.0	Male	81.0	0	0	Yes	Priv:

```
# See if data is imbalanced on the variable of interest, stroke
# Count how many '1's (stroke) and '0's (no stroke) appear
print(df_noNA["stroke"].value_counts())
# Dataset is quite unbalanced on the stroke variable
# Around 209/(4700+209) = 4.3% of dataset is positive for stroke
```

```
0    4700
1     209
Name: "stroke", dtype: int64
```



```
# Change 'stroke' attribute into data type float
df_noNA["stroke"] = df_noNA["stroke"].astype(float)
df_noNA.head(10)
```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using loc[row_index_col_index] = value instead

```

```
print(df_noNA.dtypes)
```

```

"id"          float64
"gender"      object
"age"         float64
"hypertension" object
"heart_disease" object
"ever_married" object
"work_type"   object
"residence_type" object
"avg_glucose_level" float64
"bmi"         float64
"smoking_status" object
"stroke"      float64
dtype: object

```

```

# See if there are any extreme values in numeric data
df_noNA.describe()

```

	"id"	"age"	"avg_glucose_level"	"bmi"	"stroke"
<b>count</b>	4909.000000	4909.000000	4909.000000	4909.000000	4909.000000
<b>mean</b>	37064.313506	42.865374	105.305150	28.893237	0.042575
<b>std</b>	20995.098457	22.555115	44.424341	7.854067	0.201917
<b>min</b>	77.000000	0.080000	55.120000	10.300000	0.000000
<b>25%</b>	18605.000000	25.000000	77.070000	23.500000	0.000000
<b>50%</b>	37608.000000	44.000000	91.680000	28.100000	0.000000
<b>75%</b>	55220.000000	60.000000	113.570000	33.100000	0.000000
<b>max</b>	72940.000000	82.000000	271.740000	97.600000	1.000000



```

# Normalize continuous numeric variables
# Such as age, avg_glucose_level, and bmi
# Using z-score methods

```

```

# Import libraries for normalization
from sklearn import preprocessing
scaler = preprocessing.MinMaxScaler()

```

```

# Only need to normalize continuous numeric variables
var_to_norm = ["age", "avg_glucose_level", "bmi"]
df_noNA[var_to_norm] = scaler.fit_transform(df_noNA[var_to_norm])

```

```
# Examine first 10 rows of normalized dataset
```

```
df_noNA.head()
```

```
# The 3 columns are now standardized to values between 0-1
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py:3678: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user
```

```
self[col] = igetitem(value, i)
```

	"id"	"gender"	"age"	"hypertension"	"heart_disease"	"ever_married"	"work_type"
0	9046.0	Male	0.816895	0	1	Yes	Pr
2	31112.0	Male	0.975586	0	1	Yes	Pr
3	60182.0	Female	0.597168	0	0	Yes	Pr
4	1665.0	Female	0.963379	1	0	Yes	empl
5	56669.0	Male	0.987793	0	0	Yes	Pr



```
# Create list of categorical columns
```

```
cat_cols = ["gender", "hypertension", "heart_disease", "ever_married", "work_type",
```

```
# Create copy of a data frame in memory w/ a different name
```

```
df_dummy = df_noNA.copy()
```

```
# Convert only categorical feature into dummy/one-hot features
```

```
df_dummy = pd.get_dummies(df_noNA, columns = cat_cols, prefix = cat_cols)
```

```
# Print dataset
```

```
df_dummy
```

	"id"	"age"	"avg_glucose_level"	"bmi"	"stroke"	"gender"_Female	"gender"_Male
0	9046.0	0.816895	0.801265	0.301260	1.0	0	0
2	31112.0	0.975586	0.234512	0.254296	1.0	0	0
3	60182.0	0.597168	0.536008	0.276060	1.0	1	1
4	1665.0	0.963379	0.549349	0.156930	1.0	1	1
5	56669.0	0.987793	0.605161	0.214204	1.0	0	0

```
# Create train test set split
```

```
from sklearn.model_selection import train_test_split
```

```
5106 44873 0 0 987793 0.323516 0.340206 0 0 1
```

```
# Balancing Dataset: Import libraries needed for SMOTE Oversampling method
```

```
from collections import Counter
```

```
from imblearn.over_sampling import SMOTE
```

```
import warnings
```

```
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
4909 rows x 25 columns
```

```
print(df_dummy.dtypes)
```

```
"id" float64
"age" float64
"avg_glucose_level" float64
"bmi" float64
"stroke" float64
"gender"_Female uint8
"gender"_Male uint8
"gender"_Other uint8
"hypertension"_0 uint8
"hypertension"_1 uint8
"heart_disease"_0 uint8
"heart_disease"_1 uint8
"ever_married"_No uint8
"ever_married"_Yes uint8
"work_type"_Govt_job uint8
"work_type"_Never_worked uint8
"work_type"_Private uint8
"work_type"_Self-employed uint8
"work_type"_children uint8
"residence_type"_Rural uint8
"residence_type"_Urban uint8
"smoking_status"_Unknown uint8
"smoking_status"_formerly smoked uint8
"smoking_status"_never smoked uint8
"smoking_status"_smokes uint8
dtype: object
```

```
# Partition the class of interest, stroke (Dependent Variable), from the Independent Variable
```

```
iv_classes = df_dummy.iloc[:, :-1]
```



```

dv_class = df_dummy.iloc[:, -1]

# Set class name as "stroke", all other attributes will be used as features
class_col_name = '"stroke"'
# Obtain necessary dummy feature names
dummy_feature_name = df_dummy.columns.values.tolist()
dummy_feature_names = dummy_feature_name[5:]

# 70% training, 30% test set split
x_train, x_test, y_train, y_test = train_test_split(df_dummy.loc[:, dummy_feature_names], df_

# Display the class distribution in its original split
print("Original class split prior to SMOTE oversampling: ", Counter(y_train))

    Original class split prior to SMOTE oversampling:  Counter({0.0: 3291, 1.0: 145})

# Implement the SMOTE Oversampler Model
SMOTEOverSampler = SMOTE()

# Apply data into Random Undersampler Model
x_train_SMOTE, y_train_SMOTE = SMOTEOverSampler.fit_resample(x_train, y_train)

# Display the class distribution after the SMOTE Oversampler is applied
print("Class split after SMOTE Oversampling: ",Counter(y_train_SMOTE))

    Class split after SMOTE Oversampling:  Counter({0.0: 3291, 1.0: 3291})

# Import required libraries
from sklearn.svm import SVC
from sklearn.metrics import classification_report, roc_auc_score

# Evaluate the SMOTE Oversampler
model = SVC()
clf_SMOTEOverSampler = model.fit(x_train_SMOTE, y_train_SMOTE)
pred_SMOTEOverSampler = clf_SMOTEOverSampler.predict(x_test)

# Display ROC AUC score
print("SMOTE Oversampled data - ROC AUC: ", roc_auc_score(y_test, pred_SMOTEOverSampler))

    SMOTE Oversampled data - ROC AUC:  0.5

start = time.time()

# Import needed libraries for Logistic Regression Model

```

```

from sklearn.linear_model import LogisticRegression

# Begin to Implement Logistic Regression Model
log_regr = LogisticRegression()

# Apply data into Logistic Regression Model
log_regr.fit(x_train, y_train)
y_pred = log_regr.predict(x_test)

# Obtain Confusion Matrix and Evaluation Metrics for the Logistic Regression Model
from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
cnf_matrix

array([[1409,    0],
       [  64,    0]])

# Display Evaluation Metrics for Logistic Regression Model
print("Logistic Regression Accuracy:\t", metrics.accuracy_score(y_test, y_pred))
print("Logistic Regression Precision:\t", metrics.precision_score(y_test, y_pred))
print("Logistic Regression Recall:\t", metrics.recall_score(y_test, y_pred))

Logistic Regression Accuracy:    0.956551255940258
Logistic Regression Precision:    0.0
Logistic Regression Recall:    0.0
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedWarning:
    _warn_prf(average, modifier, msg_start, len(result))

```

◀
▶

```

# Import libraries for cross-validation
from sklearn.model_selection import cross_val_score, cross_val_predict

# 10-Fold Cross Validation for Logistic Regression
cv_lr = cross_val_score(log_regr, df_dummy, df_dummy[class_col_name], cv=10)
print("Cross-validated scores:\t", cv_lr)

# Increased Accuracy score from 0.957 to 1

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: ConvergenceWarning:
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
    extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: ConvergenceWarning:
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html

```

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
Cross-validated scores: [1.          1.          1.          1.          1.          1.  
1.          1.          0.95723014 0.95918367]
```



```
# Cross validation accuracy for Logistic Regression (R2 score)  
predictions = cross_val_predict(log_regr, df_dummy, df_dummy[class_col_name], cv=10)  
accuracy = metrics.r2_score(df_dummy[class_col_name], predictions)  
print("Cross-Predicted Accuracy for Logistic Regression: ", accuracy)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: Convergence  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: Convergence  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```


Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
Cross-Predicted Accuracy for Logistic Regression: 0.7951043469408531
```



```
end = time.time()
```

```
print("Time to run Logistic Regression: ", end - start)
```

```
Time to run Logistic Regression: 3.1480841636657715
```

```
start = time.time()
```

```
# Naive Bayes modeling  
from sklearn.naive_bayes import MultinomialNB
```

```
# Create Multinomial NB Classifier  
nb = MultinomialNB()
```

```
# Train model using training sets  
nb.fit(x_train_SMOTE, y_train_SMOTE)
```

```
MultinomialNB()
```

```
# Predict response for test dataset
y_pred = nb.predict(x_test)

# Print Naive Bayes output
print("Number of features used: ", nb.n_features_)
print("Classes: ", nb.classes_)
print("Number of records for classes: ", nb.class_count_)
print("Log prior probability for classes: ", nb.class_log_prior_)
print("Log conditional probability for each feature given a class: ", nb.feature_log_prob_)

Number of features used: 20
Classes: [0. 1.]
Number of records for classes: [3291. 3291.]
Log prior probability for classes: [-0.69314718 -0.69314718]
Log conditional probability for each feature given a class: [[ -2.47270743 -2.84008951
-1.99056696 -5.08288006 -2.97245497 -2.38986067 -3.98660149
-7.15535293 -2.49664198 -3.8615758 -3.90153906 -2.67927954
-2.60089142 -3.11914766 -3.7690812 -2.91883388 -3.83311859]
[ -3.04260823 -2.72915412 -12.73907161 -3.96491332 -2.66402858
-3.69762338 -2.55636268 -2.63730714 -4.35953259 -2.21559915
-12.73907161 -2.89257778 -2.53056564 -12.73907161 -2.93711857
-3.21611268 -2.7126145 -2.46927567 -3.09325436 -2.55059311]]
```

```
# Get Naive Bayes Classifier Confusion matrix
from sklearn.metrics import confusion_matrix
cf = confusion_matrix(y_test, y_pred)
print("Confusion Matrix")
print(cf)
tn, fp, fn, tp = cf.ravel()
print("TP: ", tp, ", FP: ", fp, ", TN: ", tn, ", FN: ", fn)
```

```
Confusion Matrix
[[1357  52]
 [ 58   6]]
TP: 6 , FP: 52 , TN: 1357 , FN: 58
```

```
# Get Naive Bayes Classifier Classifier report
from sklearn.metrics import classification_report
from sklearn import metrics
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0.0	0.96	0.96	0.96	1409
1.0	0.10	0.09	0.10	64
accuracy			0.93	1473
macro avg	0.53	0.53	0.53	1473

```
weighted avg      0.92      0.93      0.92      1473
```

```
# Display Evaluation Metrics for Naive Bayes Classifier
```

```
print("Naive Bayes Classifier Accuracy:\t", metrics.accuracy_score(y_test, y_pred))
print("Naive Bayes Classifier Precision:\t", metrics.precision_score(y_test, y_pred))
print("Naive Bayes Classifier Recall:\t\t", metrics.recall_score(y_test, y_pred))
```

```
Naive Bayes Classifier Accuracy:      0.9253224711473184
Naive Bayes Classifier Precision:     0.10344827586206896
Naive Bayes Classifier Recall:        0.09375
```

```
# 10-Fold Cross Validation for Naive Bayes Classifier
```

```
cv_nb = cross_val_score(nb, df_dummy, df_dummy[class_col_name], cv=10)
print("Cross-validated scores:\t", cv_nb)
```

```
# Increased Accuracy score from 0.926 to 0.998
```

```
Cross-validated scores: [0.99796334 0.99592668 0.99796334 0.99796334 0.99592668 0.99389
0.99592668 0.99796334 0.99796334 0.99795918]
```



```
# Cross validation accuracy for Naive Bayes Classifier (R2 score)
```

```
predictions = cross_val_predict(nb, df_dummy, df_dummy[class_col_name], cv=10)
accuracy = metrics.r2_score(df_dummy[class_col_name], predictions)
print("Cross-Predicted Accuracy for Naive Bayes Classifier: ", accuracy)
```

```
Cross-Predicted Accuracy for Naive Bayes Classifier: 0.9250381757100682
```

```
end = time.time()
```

```
print("Time to run Naive Bayes Classifier: ", end - start)
```

```
Time to run Naive Bayes Classifier: 0.7576520442962646
```

```
start = time.time()
```

```
# Decision tree on dummy encoded data
```

```
from sklearn import tree
```

```
clf = tree.DecisionTreeClassifier(max_depth = 5) # 5 levels set
```

```
clf = clf.fit(x_train_SMOTE, y_train_SMOTE)
```

```
import graphviz
```

```
# Obtain unique class values to show on tree
```

```
class_values = df_dummy[class_col_name].unique()
```


```
print("class names: ", class_values)
```

```
☞ class names: [1. 0.]
```

```
# Import libraries for plotting the decision tree
```

```
import matplotlib
from matplotlib import pyplot as plt


# Plot decision tree
fig = plt.figure(figsize=(25,20))
_ = tree.plot_tree(clf, feature_names = dummy_feature_names, class_names = str(class_values),
```



```
# Save decision tree figure
fig.savefig("decision_tree8.png")
```



```
# Perform prediction on test set
y_pred = clf.predict(x_test)
```



```
# Get decision tree confusion matrix
cf = confusion_matrix(y_test, y_pred)
print("Confusion Matrix")
print(cf)
tn, fp, fn, tp = cf.ravel()
print("TP: ", tp, ", FP: ", fp, ", TN: ", tn, ", FN: ", fn)
```

```
Confusion Matrix
[[495 914]
 [ 6  58]]
TP:  58 , FP:  914 , TN:  495 , FN:  6
```

```
# Get decision tree report
from sklearn.metrics import classification_report
from sklearn import metrics
print(classification_report(y_test, y_pred))
```

```
# Oversampling can lead to overfitting
```

	precision	recall	f1-score	support
0.0	0.99	0.35	0.52	1409
1.0	0.06	0.91	0.11	64
accuracy			0.38	1473
macro avg	0.52	0.63	0.32	1473
weighted avg	0.95	0.38	0.50	1473

```
# Display Evaluation Metrics for Decision Tree
print("Decision Tree Accuracy:\t\t", metrics.accuracy_score(y_test, y_pred))
print("Decision Tree Precision:\t", metrics.precision_score(y_test, y_pred))
print("Decision Tree Recall:\t\t", metrics.recall_score(y_test, y_pred))
```

```
Decision Tree Accuracy:      0.3754243041412084
Decision Tree Precision:     0.059670781893004114
Decision Tree Recall:        0.90625
```

```
# 10-Fold Cross Validation for Decision Tree
```

```
cv_dt = cross_val_score(clf, df_dummy, df_dummy[class_col_name], cv=10)
```

```
print("Cross-validated scores:\t", cv_dt)
```

```
# Increased Accuracy score from 0.375 to 1
```

```
Cross-validated scores: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```

```
# Cross validation accuracy for Decision Tree (R2 score)
```

```
predictions = cross_val_predict(clf, df_dummy, df_dummy[class_col_name], cv=10)
```

```
accuracy = metrics.r2_score(df_dummy[class_col_name], predictions)
```

```
print("Cross-Predicted Accuracy for Decision Tree: ", accuracy)
```

```
Cross-Predicted Accuracy for Decision Tree: 1.0
```

```
end = time.time()
```

```
print("Time to run Decision Tree: ", end - start)
```

```
Time to run Decision Tree: 4.123949289321899
```